

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: `grader_weights()`, `grader_sigmoid()`, `grader_logloss()` etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

In [44]:

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

In [45]:

```
# please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant
=5,
                        n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datas
ets.make_classification.html) for more details
```

In [46]:

```
X.shape, y.shape
```

Out[46]:

```
((50000, 15), (50000,))
```

Splitting data into train and test

In [47]:

```
#please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1
5)
```

In [48]:

```
# Standardizing the data.
scaler = StandardScaler()
x_train = scaler.fit_transform(X_train)
x_test = scaler.transform(X_test)
```

In [49]:

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[49]:

```
((37500, 15), (37500,)), (12500, 15), (12500,))
```

SGD classifier

In [50]:

```
# alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15,
penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html)
```

Out[50]:

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0001,
              fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
              loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
              penalty='l2', power_t=0.5, random_state=15, shuffle=True,
              tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [51]:

```
clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.05 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.06 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.07 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.08 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.09 seconds.
Convergence after 10 epochs took 0.09 seconds
```

Out[51]:

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0001,
              fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
              loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
              penalty='l2', power_t=0.5, random_state=15, shuffle=True,
              tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [52]:

```
clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

Out[52]:

```
(array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
         0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
         0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
 (1, 15),
 array([ -0.8531383]))
```

This is formatted as code

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)

$$\begin{aligned} \logloss &= -1 \\ &\times \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} \\ &(Y_t \log_{10}(Y_{\text{pred}}) \\ &+ (1 - Y_t) \log_{10}(1 \\ &- Y_{\text{pred}})) \end{aligned}$$

- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)

$$\begin{aligned} dw^{(t)} &= x_n (y_n - \\ &\sigma((w^{(t)})^T x_n \\ &+ b^{(t)})) - \\ &\frac{\lambda}{N} w^{(t)} \end{aligned}$$

- Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this](#)

$$\begin{aligned} db^{(t)} &= y_n - \\ &\sigma((w^{(t)})^T x_n \\ &+ b^{(t)}) \end{aligned}$$

- Update weights and intercept (check the equation number 32 in the above mentioned [pdf](#)):

$$\begin{aligned} w^{(t+1)} &\leftarrow w^{(t)} + \\ &\alpha(dw^{(t)}) \\ b^{(t+1)} &\leftarrow b^{(t)} \\ &+ \alpha(db^{(t)}) \end{aligned}$$

- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
- And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training

- append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

In [53]:

```
def initialize_weights(dim):
    ''' In this function, we will initialize our weights and bias'''
    #initialize the weights to zeros array of (1,dim) dimensions
    #you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
    #initialize bias to zero
    w = np.zeros_like(dim)
    b = 0
    return w,b
```

In [54]:

```
dim=X_train[0]
w,b = initialize_weights(dim)
print('w =', (w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

In [55]:

```
dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
    assert ((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

Out[55]:

True

Compute sigmoid

$$\text{sigmoid}(z) = \frac{1}{1 + \exp(-z)}$$

In [56]:

```
#import math
def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    # compute sigmoid(z) and return
    #https://www.educative.io/edpresso/calculating-the-exponential-value-in-python
    #sigmoid_z = 1/(1+math.exp(-z))
    #https://numpy.org/doc/stable/reference/generated/numpy.exp.html
    #we can also use math library to calculate the exponential value
    sigmoid_z = 1/(1+np.exp(-z))
    return sigmoid_z
```

Grader function - 2

In [57]:

```
def grader_sigmoid(z):
    val=sigmoid(z)
    assert (val==0.8807970779778823)
```

```
        return True
grader_sigmoid(2)
```

Out[57]:

True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

In [58]:

```
def logloss(y_true, y_pred):
    '''In this function, we will compute log loss '''
    loss = 0
    n = len(y_true)
    for i in range(n):
        # we can also use either math.log10 or math.log(value,10)
        #loss += y_true[i]*math.log10(y_pred[i]) + (1-y_true[i])*math.log10(1-y_pred[i])
        #https://www.geeksforgeeks.org/numpy-log10-python/#:~:text=About%20%3A,all%20the%
        20input%20array%20elements.
        loss += y_true[i]*np.log10(y_pred[i]) + (1-y_true[i])*np.log10(1-y_pred[i])
    loss = -1 * (loss/n)

    return loss
```

Grader function - 3

In [59]:

```
def grader_logloss(true, pred):
    loss=logloss(true, pred)
    assert (loss==0.07644900402910389)
    return True
true=[1,1,0,1,0]
pred=[0.9,0.8,0.1,0.8,0.2]
grader_logloss(true, pred)
```

Out[59]:

True

Compute gradient w.r.to 'w'

$$\begin{aligned} dw^{(t)} &= x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) \\ &\quad - \frac{\lambda}{N} w^{(t)} \end{aligned}$$

In [60]:

```
def gradient_dw(x, y, w, b, alpha, N):
    '''In this function, we will compute the gardient w.r.to w '''
    #https://numpy.org/doc/stable/reference/generated/numpy.dot.html
    dw = x * (y - sigmoid(np.dot(w, x) + b)) - (alpha/N)*w
    return dw
```

Grader function - 4

In [61]:

```
def grader_dw(x, y, w, b, alpha, N):
    grad_dw=gradient_dw(x, y, w, b, alpha, N)
    #print(type(grad_dw))
    assert (np.sum(grad_dw)==2.613689585)
```

```

    return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)

```

Out[61]:

True

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

In [62]:

```

def gradient_db(x,y,w,b):
    '''In this function, we will compute gradient w.r.to b '''
    db = y - sigmoid(np.dot(w,x) + b)
    return db

```

Grader function - 5

In [63]:

```

def grader_db(x,y,w,b):
    grad_db=gradient_db(x,y,w,b)
    #print(type(grad_db))
    assert(grad_db==-0.5)
    return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b)

```

Out[63]:

True

Implementing logistic regression

In [65]:

```

def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
    ''' In this function, we will implement logistic regression'''
    train_loss = []
    test_loss = []
    w,b = initialize_weights(X_train[0])
    for i in range(0,epochs):
        #https://stackoverflow.com/questions/49783594/for-loop-and-zip-in-python
        #zip allows you to iterate two lists at the same time
        for x,y in zip(X_train,y_train):
            grad_dw = gradient_dw(x,y,w,b,alpha,N)
            grad_db = gradient_db(x,y,w,b)
            w = w + (eta0*grad_dw)
            b = b + (eta0*grad_db)

        y_train_pred = []
        for x_test in X_train:
            y_train_pred.append(sigmoid(np.dot(w,x)+b))

```

```

#computing and storing all the train loss values in a list
train_loss.append(logloss(y_train,y_train_pred))

y_test_pred = []
for x in X_test:
    y_test_pred.append(sigmoid(np.dot(w,x)+b))
#computing and storing all the test loss values in a list
test_loss.append(logloss(y_test,y_test_pred))

#https://stackoverflow.com/questions/455612/limiting-floats-to-two-decimal-points
#limiting the train and test loss upto 5 decimal places
print(f"EPOCH: {i} Train Loss: {logloss(y_train, y_train_pred):.5f} Test Loss: {logloss(y_test, y_test_pred):.5f}")

for i,loss in enumerate(train_loss):

    #print(i, loss)
    if i > 0:
        previous_loss = train_loss[i-1]
        current_loss = loss
        #print(previous_loss,current_loss)
        if(previous_loss > current_loss) and (previous_loss - current_loss < 0.001):
            break
    else:
        continue
    break

return w,b,train_loss,test_loss

```

In [66]:

```

alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=50
w,b,loss_train,loss_test=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)

```

```

EPOCH: 0 Train Loss: 0.29606 Test Loss: 0.17595
EPOCH: 1 Train Loss: 0.29128 Test Loss: 0.16940
EPOCH: 2 Train Loss: 0.29005 Test Loss: 0.16721
EPOCH: 3 Train Loss: 0.28953 Test Loss: 0.16622
EPOCH: 4 Train Loss: 0.28921 Test Loss: 0.16572
EPOCH: 5 Train Loss: 0.28899 Test Loss: 0.16546
EPOCH: 6 Train Loss: 0.28882 Test Loss: 0.16531
EPOCH: 7 Train Loss: 0.28869 Test Loss: 0.16523
EPOCH: 8 Train Loss: 0.28860 Test Loss: 0.16519

```

In [67]:

```

print(w)
print(b)
#print(train_loss)
#print(test_loss)

```

```

[-0.42101199  0.19028065 -0.14496425  0.33810803 -0.20893349  0.56366252
 -0.44548271 -0.09228306  0.21670984  0.16844415  0.1940773  0.00318758
 -0.07669313  0.33872841  0.02189547]
-0.8369332521112822

```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

In [68]:

```

# these are the results we got after we implemented sgd and found the optimal weights and intercept

```

```
w-clf.coef_, b-clf.intercept_
```

Out[68]:

```
(array([[ 0.00235492,  0.00480499,  0.00362611, -0.00333604, -0.00074678,
          0.00349673,  0.00694212,  0.00180507,  0.00743664, -0.01239711,
          -0.00297461, -0.00103157,  0.00291056,  0.0002004 , -0.00077174]]),
array([0.01620505]))
```


Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

In [69]:

```
import matplotlib.pyplot as plt

epochs = range(0,9)
plt.plot(epochs, loss_train, 'g', label='Training loss')
plt.plot(epochs, loss_test, 'b', label='Test loss')
plt.title('Training and Test loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



In [70]:

```
def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
            predict.append(1)
        else:
            predict.append(0)
    return np.array(predict)
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

0.95648

0.95416